



# IN 5380: Compiler construction

Spring 2025

## Oblig 1

18. 2. 2025

Issued: 18. 2. 2025

## 1 Official grading guidelines

The **deadline/frist** for the oblig is

**Friday, 28. 03. 2025, 23:59**

### Requirements:

- the test program parses 100% ok (see below)
- the printed AST reflects the correctly parsed structure (especially, the correct associativity and precedence). The exact same formatting as in the provided illustration is not required; reasonable deviations are fine.
- follow the instructions under “What and how to hand-in”
- the solution needs to compile and run *at the UiO pool*.<sup>1</sup> Test it!

## 2 What and how to hand in

### 2.1 Git

We use

**git**

to hand in the obligs and the two deliverables are just two “milestones” in your project. If you working in a group, git allows to share things.

I assume that most are in principle familiar with **git**, if not, ask me. We use the UiO-internal git server<sup>2</sup>

**github.uio.no**

In order to hand-in via git, each group has to do the following steps. See also the top-level readme at <https://github.uio.no/compilerconstruction-inf5110/compila/> (resp. at the corresponding place in your clone).<sup>3</sup>

<sup>1</sup>Statements like “but on my laptop it worked, ...” don’t cut it ....

<sup>2</sup>Actually, it’s the UiO enterprise github server. Github is just one particular “web-interface” on top of git which does the real work; there are also alternatives to github (and git) but for uniformity, we stick to the UiO github server.

<sup>3</sup>The name of the repository contains **infs5110**. That’s because the course was numbered 2024 and earlier *INF 5110* and the repository has not been renamed to reflect the change.

1. everyone: if not already: create yourself an **account** at `github.uio.no`. Your UiO login allows you to do that.
2. send me an email with your name (and if you have, with the name of your collaborator), so I can assign a **group number**.
3. each group will get one *repository*. For instance, group **01** get the repository **compila25-01** etc. I will add everyone to the corresponding repos as collaborator. Each repos contains the same starting point, of course.

## 2.2 What to include into a solution

- A top-level *Readme-file* containing
  - names and emails of the authors
  - easy instructions how to build the compiler and how to run it.

The top-level readme should describe in a concise manner instructions for building, installation, running, and testing your compiler in a manner useful for an interested *user*. The target audience is a master-level computer science student or someone how is not afraid to git-clone a repository (or download a jar-file or similar) and following a few clear installation steps. It should be doable “stand-alone”, i.e., without loading some files into Eclipse or similar *developer* tools (developer  $\neq$  user). The user is not expected to figure out himself or herself how to install, run, and test it. The user is also *not* expected to be a student of this compiler course or familiar with the specification of the compila-language and the oblig-documentation.

It might be nice to use some *markdown* format (like a **Readme.md**). It’s some form of poor-man’s markup, and is typically rendered nicely by browsers. The current top-level Compila-Readme is not in the **md**-format, but in **org**-format (**Readme.org**) (which is a similar, but a more expressive markdown format and also rendered nicely by github, and which personally I prefer for myself).

- The information discussed so far is for an “interested user” which may be an outsider, maybe like a random surfer on the internet who wants to try it out. Besides that, **additional** information should be provided, which is not relevant for an outsider, but contains “insider information” relevant for the course’s mandatory assignment and the technical realization of the solution. Thus, that information might perhaps better not be kept in the top level readme but at a different, less prominent place or file.
  - test-output for running the compiler on `compila.cmp` as input
  - of course, the *code* needed to run your package. That includes
    - \* *JFlex*-code for the scanner
    - \* *CUP*-code for the 2 variants of the syntax
    - \* the Java-classes for the syntax-tree
    - \* the build-script `build.xml`<sup>4</sup>

## 2.3 Delivery in a nutshell

Ideal would be an

### out-of-the-box

behavior. Assuming that the “outside user” has the required software in place (like JDK and *ant*, or something rather common, and also is familiar with git). Then it works out-of-the-box, by following simple steps like for example (for the repos `compila25-<n>` where *n* is the group number):

```
git clone <repos-url>
cd compila25-<n>
ant
ant build
ant run
```

and those steps (whatever they are) should be mentioned clearly in your Readme. The steps mentioned here correspond to those suggested when using the provided starting point, which corresponds to

<sup>4</sup>Alternatively, you can use a makefile if you prefer that. In that case, remove the `build.xml`-file.

<https://github.uio.no/compilerconstruction-inf5110/compila>.<sup>5</sup>

Your solution may involve different steps (maybe **ant** **ast** for producing the require AST), and may require adaptation of `build.xml` (if you use ant). But you should **keep the spirit** of “out-of-the-box” and “kiss” (keep it simple, stupid.). Remember: the user is always stupid...

### 3 Purpose and goal

The goal of the task is to gather practical experience of the following tools and techniques.

- use scanner/lexer and parser tools. In this case **JFlex** and **CUP**.
- rewrite and massage a grammar given in one form into another one so that it's accepted by the tools. In our case, the language is given in some EBNF, which has to be adequately rewritten so that it can be fit into the lexer and parser tools
- handle associativity and precedence of syntactic constructs in **one** of two possible ways
  - formulate a (unambiguous) grammar that embodies the correct precendences and associativities
  - work with an ambiguous grammar, but instruct the parser tool (like CUP) to result in an appropriate parser.
- design and implement an suitable AST data structure. Use the parser to output your AST (in case of a successful parse).
- do a “*pretty printer*” in the following sense: implement some functionality that *outputs* and AST in a “useful” manner. In particular, the parentetic tree structure must be visible from the output (i.e., one can see whether the associativity and the precedence is correctly implemented).

### 4 Default tools

The platform is *Java*, together with the auxiliary tools

- *JFlex* (scanner generator in the (f)lex family)
- *CUP* (parser generator in the yacc family)
- *ant* (a kind of “make” tool specialized for Java)

The tool ant is available at the RHEL pool at IFI. *JFlex* and *CUP* are provided as part of the repos.

If, for some reason, you plan to deviate from the suggested tools, you

1. **MUST** discuss that first with the lecturer.
2. it **must** be a platform which is freely available at the university RHEL pool resp. is platform independent. Proprietary tools, tools for which one has to pay, or tools I don't have easy<sup>6</sup> access to cannot be used. If using Java, it must compile and run without support of specific development environments or “frameworks” besides the ones mentioned (*JFlex*, *CUP*, *ant*).

### 5 First steps (“oblig 0”)

Apart from getting yourself a group number, git access etc. as described, there are things you really should do afterwards:

#### Checking that on your platform, all is in place

<sup>5</sup>Don't clone this one for doing the oblig, clone the repos corresponding to your group. If you clone it, it's the same information, but you have no write access to the repos.

<sup>6</sup>I *mean* easy and in the sense that it does not cost money or much time to install the required environment or to figure out how it all hangs together. *Not* “easy” as in “it's really not so hard after you read some manuals and played around with it for a while and with the help from the fine folks on stack-exchange” ...

The provided readme describes the easy steps (not repeated here). Make sure very early on that your set-up works, that ant is running, etc. If the tests fail, **contact me to get that out of the way** before the actual oblig work starts. The starting point provides those tests to avoid that later and shorter before the deadline, you find out “something’s not working on my machine”. I tested it on the university machines, I tested it on my laptop, but I cannot test it on other people’s machines. Normally it’s not a problem, the stuff is supposed to be platform independent, but “normally ok” means, sometimes, maybe only rarely, it’s not ok, and sometimes for weird reasons.

To avoid running into frustrating stupid things later is **only one purpose** to run the provided the tests and make sure all is in order. The **other purpose** is: the tests run lexers and a parsers! It’s of course not the compila language, that’s your task, but it shows in two tiny examples, how the tools operate together, and how a small lex file looks and how a small cup-file looks etc. And for that:

**Do that checks immediately** even if you don’t plan to start on the content your oblig right away.

It’s a check that all is in place.

Talking about how to tackle practical aspects of the oblig(s) and organizing the task. You may read the advice I typed up a couple of years ago about that

<https://martinsteffen.github.io/compilerconstruction/howtowriteaparser/>

## 6 Task more specifically: Syntax check and parsing

The overall task is to

implement a parser for the *Compila 25* language.

The language specification is given in a separate document. Oblig 1 is concerned with checking *syntactic correctness*, which means, not all of the language specification is relevant right now: semantic correctness, type checking etc. will become relevant only later for the second oblig.

### 6.1 Syntax tree

The result of a successful parse is an *abstract syntax tree*. That data structure needs to be appropriately “designed”. In a Java implementation, that involves the definition of appropriately chosen classes arranged in some class hierarchy. Make also use if *abstract classes*. In the lecture, there had been some “design guidelines” that may be helpful. Carefully chosen names for classes will help in a conceptually clear implementation. A definitely *non-recommended* way is to have one single class `Node` lumping together all kinds of nodes and syntactic categories in the syntax tree.

### 6.2 Print out of the AST

The AST should be “printed”. The easiest and recommended form of printout is in *prefix form*. Under `material/sample-compila-ast`, there is an example compila input file and a corresponding file containing a possible output. The two files are called

- `complexaddition.cmp`
- `complexaddition.ast`

Note: the two files are meant as *inspiration*. Each year the syntax of *compila* slightly changes (wrt. keywords, associativity etc). So the syntax is **not 100%** in accordance with the 2025 version (but pretty comparable).

The one that **should actually be used** for this task (AST-printing) is the following:

```
./src/tests/fullprograms/complexaddition.cmp
```

and this one should be consistent with this year’s grammar (fingers crossed).

It’s allowed (but not necessary) to print it in other forms than the prefix form used in the inspiration. But the output must indicate the AST in readable form (“readable” as in human-readable that is ...). Note, the task is not that the output is a syntactically correct *compila* program again (that might be a so-called formatting tool or a pretty-printer), we just need a way to look at the syntax tree, which comes in handy for debugging,

### 6.3 One of two grammars

As mentioned shortly, one can approach some ambiguities in the grammar, the ones in connection with precedence and associativity, in one of two ways:

1. formulating an *unambiguous* grammar resolving precedence and associativity by “baking it in” directly into the grammar (“precedence cascade”). The grammar is in plain BNF (in the form required by the tools)
2. Or one uses an ambiguous grammar and relies on *CUP* to resolve the associativity and precedence.

### Comparison and discussion

Investigate and characterize *conflicts* of the *original* grammar. How many states do the 2 generated CUP grammars have? That requires a look into the CUP-generated code. Discuss also whether the choice of the two grammars influences the generation of the AST: is one of the two approaches easier to work with when it comes to generate an AST (resp. your chosen AST data structure).

**Note:** It's not required to provide code to build *two* versions of AST-generation, one is enough. In other words, for one of the two grammars, you don't need “action code” in the grammar to produce an AST, plain *checking* is sufficient.

### 6.4 Lexical analysis

As mentioned, *JFlex* is the tool of choice for lexical analysis. It delivers a token to the parser via the method `next_token()`.

As far the the “conceptual” task concerning *compila 25* is concerned, the lexer is responsible for ignoring comment, white-space etc, find keywords and the like.

Practically, one has to make the parser and the lexer “work together” hand in hand. Information about that can be found in the corresponding manuals. There should also be examples for inspiration. A crucial ingredient is the interface `java.cup.runtime.Scanner` which needs to be implemented by the actual scanner. The scanner will hand over tokens of the type `Symbol` and one can use `Symbol.value` to pass “text” or other objects from the lexer to the parser.

### 6.5 Error handling

Error handling can be done simple: When hitting an error, parsing should stop (as opposed to try to continue and give back an avalanche of subsequent errors). Some meaningful error message (at least wrt. which syntactic class caused the error) would be welcome, as opposed to a plain “sorry, bad program”. It's not required to give back line numbers referring to the original source code or positions in the original file. In practice that's definitely useful (and not very hard either), but not required for the oblig.

### 6.6 Tests

For testing, there is a bunch of files under

```
./src/tests
```

They are supposed to contain syntactically correct programs for this year's version, with exception of the ones under

```
./src/tests/errors
```

Later for **oblig 2** there will be more tests, especially for testing type checking and code generation.

For **oblig 1**, you are requested to generate an ast for

```
./src/tests/fullprograms/complexadditions.cmp
```

as part of the oblig, as mentioned, but you may of course use the other test program as well to see how robust your implementation is.

## 7 Resources

The web-page (the git-one) will contain also links to *JFlex* *CUP* and corresponding manuals.